

facebook

Artificial Intelligence Research

How to model an RL problem: Markov Decision Processes

Matteo Pirotta

Facebook AI Research

Acknowledgments

Special thanks to Alessandro Lazaric for providing these slides from the RL class we teach in Paris.


Outline

1 Markov Decision Process

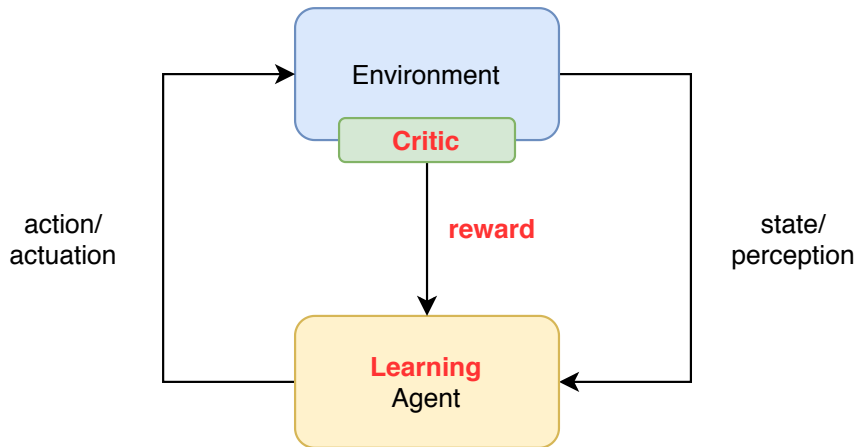
2 Policy

3 Optimality Principle

The Reinforcement Learning Model



The Reinforcement Learning Model



The RL interaction

In each *discrete* decision time $t = 1, 2, \dots$, the learning agent

- selects an action a_t based on the current state s_t (or possibly all previous observations)
- observes a reward r_t
- moves to a new state s_{t+1}

Markov Chains

Definition (Markov chain)


Let the *state space* S be a bounded compact subset of the Euclidean space, the discrete-time dynamic system $(s_t)_{t \in \mathbb{N}} \in X$ is a Markov chain if it satisfies the *Markov property*

$$\mathbb{P}(s_{t+1} = s \mid s_t, s_{t-1}, \dots, s_0) = \mathbb{P}(s_{t+1} = s \mid s_t),$$

Given an initial state $s_0 \in S$, a Markov chain is defined by the *transition probability* p

$$p(s' | s) = \mathbb{P}(s_{t+1} = s' | s_t = s).$$

Markov Chain



4 states Markov chain

Markov Decision Process

Definition (Markov decision process [1, 4, 3, 6, 2])

A **Markov decision process** is defined as a tuple $M = (S, A, p, r)$ where

- S is the *state* space,
- A is the *action* space,
- $p(s'|s, a)$ is the *transition probability* with

$$p(s'|s, a) = \mathbb{P}(s_{t+1} = s' | s_t = s, a_t = a),$$

- $r(s, a, s')$ is the *reward* of transition (s, a, s') .

Markov Decision Process

Definition (Markov decision process [1, 4, 3, 6, 2])

A **Markov decision process** is defined as a tuple $M = (S, A, p, r)$ where

- S is the **state** space,
 - A is the **action** space,
 - $p(s'|s, a)$ is the **transition probability** with
- } **finite**

$$p(s'|s, a) = \mathbb{P}(s_{t+1} = s' | s_t = s, a_t = a),$$

- $r(s, a, s')$ is the **reward** of transition (s, a, s') .
- } **bounded and simplified to $r(s, a)$**

Markov Decision Process

Definition (Markov decision process [1, 4, 3, 6, 2])

A **Markov decision process** is defined as a tuple $M = (S, A, p, r)$ where

- S is the **state** space,
- A is the **action** space,
- $p(s'|s, a)$ is the **transition probability** with

$$p(s'|s, a) = \mathbb{P}(s_{t+1} = s' | s_t = s, a_t = a),$$

- $r(s, a, s')$ is the **reward** of transition (s, a, s') . } **bounded and simplified to $r(s, a)$**

👍 Reward may be stochastic

- $\nu(s, a)$ is the reward distribution for (s, a) and

$$r(s, a) = \mathbb{E}_{R \sim \nu(s, a)}[R]$$

Markov Decision Process

Definition (Markov decision process [1, 4, 3, 6, 2])

A **Markov decision process** is defined as a tuple $M = (S, A, p, r)$ where


- S is the **state** space,
- A is the **action** space,
- $p(s'|s, a)$ is the **transition probability** with

$$p(s'|s, a) = \mathbb{P}(s_{t+1} = s' | s_t = s, a_t = a),$$

- $r(s, a, s')$ is the **reward** of transition (s, a, s') .

☞ The process generates **trajectories** $h_t = (s_1, a_1, \dots, s_{t-1}, a_{t-1}, s_t)$, with $s_{t+1} \sim p(\cdot | s_t, a_t)$

Example: Tabular MDP



- $\mathcal{S} = \{s_0, s_1, s_2\}$
- $\mathcal{A} = \{a_0, a_1\}$ ($\mathcal{A}_{s_0} = \{a_0\}$, $\mathcal{A}_{s_1} = \{a_0\}$, $\mathcal{A}_{s_2} = \{a_0, a_1\}$)
- Mean reward in s_1

$$r(s_1, a_0) = 2/3 \quad r(s_1, a_1) = 0$$

- Transition probabilities in s_0 by taking action a_0

$$p(s_1|s_0, a_0) = 0.1 \quad p(s_2|s_0, a_0) = 0.9$$

Markov Decision Process: the Assumptions



Markov assumption: the current state s and action a are a sufficient statistics for the next state s'

$$p(s'|s, a) = \mathbb{P}(s_{t+1} = s' | s_t = s, a_t = a)$$

Possible relaxations

- Possible to extend to continuous state-action space
- Define a new state $x_t = (s_t, s_{t-1}, s_{t-2}, \dots)$ (i.e., k -order MDP)
- Move to partially observable MDP (PO-MDP)
- Move to predictive state representation (PSR) model

ATARI Breakout



* figure from [5]

ATARI Breakout


$$\mathbb{P} \left[s_{t+1} = \text{[Screenshot 1]} \mid s_t = \text{[Screenshot 2]}, \text{no-move} \right]$$

Non-Markov dynamics

* figure from [5]

ATARI Breakout

12



4 consecutive frames = single observation

Markov Decision Process: the Assumptions

Time assumption: time is discrete

$$t \rightarrow t + 1$$

Possible relaxations

- Identify the proper time granularity
- Most of MDP literature extends to continuous time

ATARI Breakout

$a_t = \text{left}$



t

ATARI Breakout

$a_t = \text{left}$



$t + 1$

Too fine-grained resolution

ATARI Breakout

$a_t = \text{left}$



t

ATARI Breakout

$a_t = \text{left}$



$t + 1$

Too coarse-grained resolution

Markov Decision Process: the Assumptions

Reward assumption: the reward is uniquely defined by a transition (or part of it)

$$r(x, a, y)$$

Possible relaxations

- Distinguish between global goal and reward function
- Move to inverse reinforcement learning (IRL) to induce the reward function from desired behaviors

Markov Decision Process: the Assumptions

Stationarity assumption: the dynamics and reward do not change over time

$$p(y|x, a) = \mathbb{P}(x_{t+1} = y | x_t = x, a_t = a) \quad r(x, a, y)$$

Possible relaxations

- Identify and remove the non-stationary components (e.g., cyclo-stationary dynamics)
- Identify the time-scale of the changes
- Work on finite horizon problems

Example: the Retail Store Management Problem

Description. At each month t , a store contains s_t items of a specific goods and the demand for that goods is D_t . At the end of each month the manager of the store can order a_t more items from the supplier. Furthermore we know that

- The *cost* of maintaining an inventory of s is $h(s)$.
- The *cost* to order a items is $C(a)$.
- The *income* for selling q items is $f(q)$.
- If the demand D is bigger than the available inventory s , customers that cannot be served leave.
- The *value of the remaining inventory* at the end of the year is $g(s)$.
- *Constraint*: the store has a maximum capacity M .
- *Goal*: maximize some measure of profit

Example: the Retail Store Management Problem

- *State space*: $s \in S = \{0, 1, \dots, M\}$.

Example: the Retail Store Management Problem

- *State space*: $s \in S = \{0, 1, \dots, M\}$.
- *Action space*: it is not possible to order more items than the capacity of the store, then the action space should depend on the current state. Formally, at state s ,
 $a \in A(s) = \{0, 1, \dots, M - s\}$.

Example: the Retail Store Management Problem

- *State space*: $s \in S = \{0, 1, \dots, M\}$.
- *Action space*: it is not possible to order more items than the capacity of the store, then the action space should depend on the current state. Formally, at state s , $a \in A(s) = \{0, 1, \dots, M - s\}$.
- *Dynamics*: $s_{t+1} = [s_t + a_t - D_t]^+$.
Problem: the dynamics should be Markov and stationary!

Example: the Retail Store Management Problem


- **State space:** $s \in S = \{0, 1, \dots, M\}$.
- **Action space:** it is not possible to order more items than the capacity of the store, then the action space should depend on the current state. Formally, at state s , $a \in A(s) = \{0, 1, \dots, M - s\}$.
- **Dynamics:** $s_{t+1} = [s_t + a_t - D_t]^+$.
Problem: the dynamics should be Markov and stationary!
- The demand D_t is *stochastic and time-independent*. Formally, $D_t \stackrel{i.i.d.}{\sim} \mathcal{D}$.

Example: the Retail Store Management Problem

- **State space:** $s \in S = \{0, 1, \dots, M\}$.
- **Action space:** it is not possible to order more items than the capacity of the store, then the action space should depend on the current state. Formally, at state s , $a \in A(s) = \{0, 1, \dots, M - s\}$.
- **Dynamics:** $s_{t+1} = [s_t + a_t - D_t]^+$.
Problem: the dynamics should be Markov and stationary!
- The demand D_t is **stochastic and time-independent**. Formally, $D_t \stackrel{i.i.d.}{\sim} \mathcal{D}$.
- **Reward:** $r_t = -C(a_t) - h(s_t + a_t) + f([s_t + a_t - s_{t+1}]^+)$.

Exercise: the Parking Problem

A driver wants to park his car as close as possible to the restaurant.



- The driver cannot see whether a place is available unless he is in front of it.
- There are P places.
- At each place i the driver can either move to the next place or park (if the place is available).
- The closer to the restaurant the parking, the higher the satisfaction.
- If the driver doesn't park anywhere, then he/she leaves the restaurant and has to find another one.

1 Markov Decision Process

2 Policy

3 Optimality Principle


Policy

Definition (Policy)


A *decision rule* d can be

- *Deterministic*: $d : S \rightarrow A$,
- *Stochastic*: $d : S \rightarrow \Delta(A)$,
- *History-dependent*: $d : H_t \rightarrow A$,
- *Markov*: $d : S \rightarrow \Delta(A)$,

Probability distribution
over actions



Space of histories
 $h_t = (s_1, a_1, s_2, \dots, s_t)$



Policy

Definition (Policy)

A *decision rule* d can be

- *Deterministic*: $d : S \rightarrow A$,
- *Stochastic*: $d : S \rightarrow \Delta(A)$,
- *History-dependent*: $d : H_t \rightarrow A$,
- *Markov*: $d : S \rightarrow \Delta(A)$,

Probability distribution
over actions

Space of histories
 $h_t = (s_1, a_1, s_2, \dots, s_t)$

A *policy* (strategy, plan) can be

- *Non-stationary*: $\pi = (d_0, d_1, d_2, \dots)$,
- *Stationary*: $\pi = (d, d, d, \dots)$.

👉 An agent behaving under policy π selects at round t the action

$$a_t \sim d_t(s_t)$$

Example: the Retail Store Management Problem

- Stationary policy composed of deterministic Markov decision rules

$$\pi(s) = \begin{cases} M - s & \text{if } s < M/4 \\ 0 & \text{otherwise} \end{cases}$$

- Stationary policy composed of stochastic history-dependent decision rules

$$\pi(s_t) = \begin{cases} \mathcal{U}(M - s, M - s + 10) & \text{if } s_t < s_{t-1}/2 \\ 0 & \text{otherwise} \end{cases}$$

- Non-stationary policy composed of deterministic Markov decision rules

$$d_t(s) = \begin{cases} M - s & \text{if } t < 6 \\ \lfloor (M - s)/5 \rfloor & \text{otherwise} \end{cases}$$

Markov Chain of a Policy


Under a stationary policy $\pi : \mathcal{S} \rightarrow \Delta(\mathcal{A})$, the random process $(s_t)_{t \in \mathbb{N}}$ is a *Markov Chain*, with transition probability

$$P^\pi(s'|s) = \mathbb{P}(s_{t+1} = s' | s_t = s, \pi) = \sum_{a \in \mathcal{A}} \pi(s, a) p(s' | s, a)$$

Example: only 2 deterministic stationary policies

$$\pi_0 = \{a_0, a_0, a_0\}$$

$$\pi_1 = \{a_0, a_0, a_1\}$$



Markov Chain of a Policy


Under a stationary policy $\pi : \mathcal{S} \rightarrow \Delta(\mathcal{A})$, the random process $(s_t)_{t \in \mathbb{N}}$ is a *Markov Chain*, with transition probability

$$P^\pi(s'|s) = \mathbb{P}(s_{t+1} = s' | s_t = s, \pi) = \sum_{a \in \mathcal{A}} \pi(s, a) p(s' | s, a)$$

Example: only 2 deterministic stationary policies

$$\pi_0 = \{a_0, a_0, a_0\}$$

$$\pi_1 = \{a_0, a_0, a_1\}$$



Markov Chain of a Policy


Under a stationary policy $\pi : \mathcal{S} \rightarrow \Delta(\mathcal{A})$, the random process $(s_t)_{t \in \mathbb{N}}$ is a *Markov Chain*, with transition probability

$$P^\pi(s'|s) = \mathbb{P}(s_{t+1} = s' | s_t = s, \pi) = \sum_{a \in \mathcal{A}} \pi(s, a) p(s' | s, a)$$

Example: only 2 deterministic stationary policies

$$\pi_0 = \{a_0, a_0, a_0\}$$

$$\pi_1 = \{a_0, a_0, a_1\}$$



👍 A MDP is sometimes referred as controlled Markov chain

1 Markov Decision Process

2 Policy

3 Optimality Principle

What is the “utility” of a policy?
i.e., how good is a policy

State Value Function

Given a policy $\pi = (d_1, d_2, \dots)$ (deterministic to simplify notation)

- *Finite time horizon* T : deadline at time T , the agent focuses on the sum of the rewards up to T .

$$V^\pi(t, s) = \mathbb{E} \left[\sum_{\tau=t}^{T-1} r(s_\tau, d_\tau(h_\tau)) + R(s_T) \mid s_t = s; \pi = (d_1, \dots, d_T) \right],$$

where R is a value function for the final state.

State Value Function

Given a policy $\pi = (d_1, d_2, \dots)$ (deterministic to simplify notation)

- *Finite time horizon T* : deadline at time T , the agent focuses on the sum of the rewards up to T .

$$V^\pi(t, s) = \mathbb{E} \left[\sum_{\tau=t}^{T-1} r(s_\tau, d_\tau(h_\tau)) + R(s_T) \mid s_t = s; \pi = (d_1, \dots, d_T) \right],$$

where R is a value function for the final state.

- *Used when*: there is an intrinsic deadline to meet.

State Value Function

Given a policy $\pi = (d_1, d_2, \dots)$ (deterministic to simplify notation)

- *Infinite time horizon with discount*: the problem never terminates but rewards which are *closer* in time receive a *higher* importance.

$$V^\pi(s) = \mathbb{E} \left[\sum_{t=0}^{\infty} \gamma^t r(s_t, d_t(h_t)) \mid s_0 = s; \pi \right],$$

with discount factor $0 \leq \gamma < 1$:

- *small* = short-term rewards, *big* = long-term rewards
- for any $\gamma \in [0, 1)$ the series always converge (for bounded rewards)

State Value Function

Given a policy $\pi = (d_1, d_2, \dots)$ (deterministic to simplify notation)

- *Infinite time horizon with discount*: the problem never terminates but rewards which are *closer* in time receive a *higher* importance.

$$V^\pi(s) = \mathbb{E} \left[\sum_{t=0}^{\infty} \gamma^t r(s_t, d_t(h_t)) \mid s_0 = s; \pi \right],$$

with discount factor $0 \leq \gamma < 1$:

- *small* = short-term rewards, *big* = long-term rewards
- for any $\gamma \in [0, 1)$ the series always converge (for bounded rewards)
- *Used when*: there is uncertainty about the deadline and/or an intrinsic definition of discount.

State Value Function

Given a policy $\pi = (d_1, d_2, \dots)$ (deterministic to simplify notation)

- *Stochastic shortest path*: the problem never terminates but the agent will eventually reach a *termination state*.

$$V^\pi(s) = \mathbb{E} \left[\sum_{t=0}^{T_\pi} r(s_t, d_t(h_t)) \mid s_0 = s; \pi \right],$$

where T_π is the first (*random*) time when the *termination state* is achieved.

State Value Function

Given a policy $\pi = (d_1, d_2, \dots)$ (deterministic to simplify notation)

- *Stochastic shortest path*: the problem never terminates but the agent will eventually reach a *termination state*.

$$V^\pi(s) = \mathbb{E} \left[\sum_{t=0}^{T_\pi} r(s_t, d_t(h_t)) | s_0 = s; \pi \right],$$

where T_π is the first (*random*) time when the *termination state* is achieved.

- *Used when*: there is a specific goal condition.

State Value Function

Given a policy $\pi = (d_1, d_2, \dots,)$ (deterministic to simplify notation)

- *Infinite time horizon with average reward*: the problem never terminates but the agent only focuses on the (expected) *average of the rewards*.

$$\rho^\pi(s) = \lim_{T \rightarrow \infty} \mathbb{E} \left[\frac{1}{T} \sum_{t=0}^{T-1} r(s_t, d_t(h_t)) \mid s_0 = s; \pi \right].$$

State Value Function

Given a policy $\pi = (d_1, d_2, \dots,)$ (deterministic to simplify notation)

- *Infinite time horizon with average reward*: the problem never terminates but the agent only focuses on the (expected) *average of the rewards*.

$$\rho^\pi(s) = \lim_{T \rightarrow \infty} \mathbb{E} \left[\frac{1}{T} \sum_{t=0}^{T-1} r(s_t, d_t(h_t)) \mid s_0 = s; \pi \right].$$

- *Used when*: the system should be constantly controlled over time.

State Value Function

Technical note: the expectations refer to all possible stochastic trajectories.
A (possibly non-stationary stochastic) policy π applied from state s_0 returns

$$(s_0, r_0, s_1, r_1, s_2, r_2, \dots)$$

where $r_t = r(s_t, d_t(h_t))$ and $s_t \sim p(\cdot | s_{t-1}, a_{t-1} = d_{t-1}(h_{t-1}))$ are *random* realizations.

The value function (discounted infinite horizon) is

$$V^\pi(s) = \mathbb{E}_{(s_1, s_2, \dots)} \left[\sum_{t=0}^{\infty} \gamma^t r(s_t, d_t(h_t)) \mid s_0 = s; \pi \right]$$

Example: the Retail Store Management Problem

Simulation

Optimization Problem

Definition (Optimal policy and optimal value function)

The solution to an MDP is an *optimal policy* π^* satisfying

$$\pi^* \in \arg \max_{\pi \in \Pi} V^\pi$$

in all the states $x \in X$, where Π is some policy set of interest.

Optimization Problem

Definition (Optimal policy and optimal value function)

The solution to an MDP is an *optimal policy* π^* satisfying

$$\pi^* \in \arg \max_{\pi \in \Pi} V^\pi$$

in all the states $x \in X$, where Π is some policy set of interest.

The corresponding value function is the *optimal value function*

$$V^* = V^{\pi^*}$$

Optimization Problem

Preview of next chapter

- 1 $\pi^* \in \arg \max(\cdot)$ and not $\pi^* = \arg \max(\cdot)$ because an MDP may admit **more than one** optimal policy
- 2 π^* achieves the largest possible value function in **every** state
- 3 there always exists an optimal **deterministic** policy
- 4 except for finite-horizon problems, there always exists an optimal **stationary** policy
- 5 there exist **efficient** algorithms to compute value function and optimal policies

Limitations: Average Case

- 1 All the previous value functions define an objective **in expectation**
- 2 Other **utility functions** may be used
- 3 Risk measures could be integrated but they may induce “weird” problems and make the solution more difficult

Summary

- 1 Definition of the Markov Decision Process and its assumptions
- 2 Definition of a policy
- 3 Definition of value functions

Bibliography

- [1] R. E. Bellman. *Dynamic Programming*. Princeton University Press, Princeton, N.J., 1957.
- [2] D.P. Bertsekas and J. Tsitsiklis. *Neuro-Dynamic Programming*. Athena Scientific, Belmont, MA, 1996.
- [3] W. Fleming and R. Rishel. *Deterministic and stochastic optimal control*. Applications of Mathematics, 1, Springer-Verlag, Berlin New York, 1975.
- [4] R. A. Howard. *Dynamic Programming and Markov Processes*. MIT Press, Cambridge, MA, 1960.
- [5] Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Alex Graves, Ioannis Antonoglou, Daan Wierstra, and Martin A. Riedmiller. Playing atari with deep reinforcement learning. *CoRR*, abs/1312.5602, 2013.
- [6] M.L. Puterman. *Markov Decision Processes Discrete Stochastic Dynamic Programming*. John Wiley & Sons, Inc., New York, Etats-Unis, 1994.



Thank you!

facebook

Artificial Intelligence Research



. \ |